

ADVANCED DATA ANALYSIS FOR PSYCHOLOGICAL SCIENCE

Extra slides: Introduction to R

Luca Menghini Ph.D.

luca.menghini@unipd.it

Master degree in Developmental and Educational Psychology

University of Padova

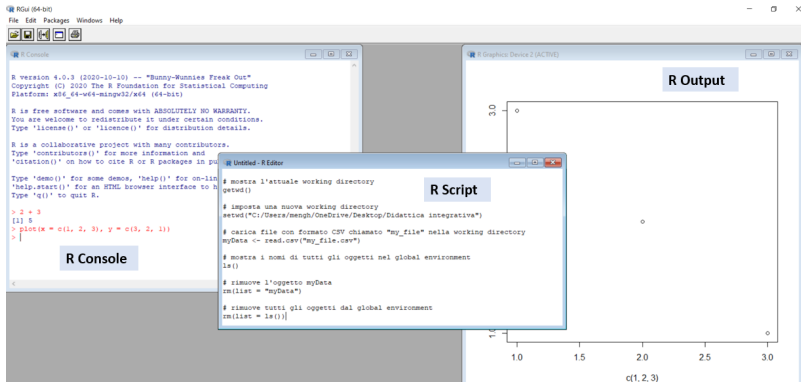
2023-2024



How to download and install

1. Download R from <https://www.r-project.org/>
CRAN (*Comprehensive R Archive Network*): network of servers that provide upated software versions and related documentation
2. Select **CRAN** from the *Download* menu on the left, then select a mirror mirror (e.g., [the first](#) or [the uniPD' one](#)), then select your OS (Linux, MacOS, or Windows)
3. Install R by opening the downloaded **.exe** (Windows) or **.pkg** (MacOS) or by following the commands required by your Linux version

The basic R interfance



- **Console:** to rapidly write (`>`) and execute (`Enter` key) the commands
- **Script** (File menu > New R Script): to write, modify, and save sequences of commands (saved with the `.R` format)
- **Outputs** (e.g., plot): graphical windows that opens when you run the command

RStudio

RStudio is a development environment for R using an **optimized graphical interface** to make it more simple (e.g., accessing files, objects, plots, datasets, etc.).

It was founded in 2009 by J. J. Allaire (written in Java and C++) and is managed and developed by an international research group. Like R, it is **free and open-source** (GNU General Public Licence) + paid premium versions.

How to install RStudio

(note: *after* installing R)

1. Download RStudio from <https://rstudio.com>
2. Select **Download** in the top menu, then select the “free” version of **RStudio Desktop**, then select your OS
3. Install RStudio by opening the downloaded file

The RStudio interface

The screenshot displays the RStudio interface with the following components:

- Source Pane:** Contains R code for creating a dataset and plotting a boxplot.

```
1 sleep # stampo dataset sleep
2
3
4 plot(extra-group,data=sleep) # boxplot
5
6
7
```
- Environment Pane:** Shows the current R objects in the workspace.

R Objects	
ab	36 obs. of 3 variables
abilities	36 obs. of 3 variables
att	210 obs. of 2 variables
- Console Pane:** Shows the output of the executed R code.

```
8 0.8 1 8
9 0.0 1 9
10 2.0 1 10
11 1.9 2 1
12 0.8 2 2
13 1.1 2 3
14 0.1 2 4
15 -0.1 2 5
16 4.4 2 6
17 5.5 2 7
18 1.6 2 8
19 4.6 2 9
20 3.4 2 10
> plot(extra-group,data=sleep) # grafico
> |
```
- Plots Pane:** Displays a boxplot of the 'extra' variable across two groups (1 and 2). The y-axis is labeled 'extra' and ranges from -1 to 5. The x-axis is labeled 'group'.

- **Source:** R Scripts (.R), documents and presentations (.Rmd), applications (.app), etc. To run one or multiple commands, select them and tap on **Ctrl + Enter** or click on the **Run** button on the top-right of the window
- **Environment:** objects included in the workspace and **History** (history of executed commands)

Some elementary commands

Comments (#)

```
# this is a comment
```

Simple mathematical operations

```
2 + 2 # sum
```

```
[1] 4
```

```
2 * 2 # multiplication
```

```
[1] 4
```

```
log(3) # natural logarithm
```

```
[1] 1.098612
```

```
exp(1) # exponential function
```

```
[1] 2.718282
```

Longer expressions (with **round brackets**)

```
sqrt(5) * ( (4 - 1/2)^2 - pi/2^(1/3) )
```

```
[1] 21.81623
```

Assigning values to **objects** (<-)

```
x <- 3 # creates the 'x' object with value 3
```

```
x # print the value of x
```

```
[1] 3
```

Objects' names can include letters, numbers, underscores, and, dots (e.g. `tony`, `tony32`, `tony.32`, `tony_32`)

```
tony_32 <- x / 3 # assign a value to the object
```

```
tony_32 # print the object's value
```

```
[1] 1
```

R is key-sensitive!

But it is not sensitive to spaces

```
3+2
```

```
[1] 5
```

```
3 + 2
```

```
[1] 5
```

Hands on : Arithmetic operations

Perform the following operations using R. ([solutions](#)):

Source: [Psicostat intro to R](#)

1. $\frac{(45+21)^3 + \frac{3}{4}}{\sqrt{32 - \frac{12}{17}}}$
2. $\frac{\sqrt{7-\pi}}{3(45-34)}$
3. $\sqrt[3]{12 - e^2} + \ln(10\pi)$
4. $\frac{\sin(\frac{3}{4}\pi)^2 + \cos(\frac{3}{2}\pi)}{\log_7 e^{\frac{3}{2}}}$
5. $\frac{\sum_{n=1}^{10} n}{10}$

Extra: Assign the result of operation #4 to the object **x**, then assign the result of operation #3 to the object **y**, and sum **x** and **y**. The result should be 5.76.

Hands on : Relational & logical operators

Relational operators

```
3 == 3 # equal to
[1] TRUE

3 != 3 # different from
[1] FALSE

x >= 3 # higher than or equal to
[1] TRUE

5 %in% c(3, 5, 8) # included in
[1] TRUE
```

Logical operators

```
x <- TRUE
y <- !x # negation of
y
[1] FALSE

x & (5 < 2) # joined condition
[1] FALSE

x | (5 < 2) # inclusive disjoint
[1] TRUE
```

Exercises on relational & logical operators:

Source: [Psicostat intro to R](#)

1. Define a proposition to test the following condition: “ x is a number included between -4 and -2 or a number included between 2 and 4 ”
2. Define two TRUE relationships and two FALSE relationship that let you test the results of all possible combinations using the logical operators & and |
3. Run the following operations: $4 \wedge 3 \%in\% c(2,3,4)$ and $4 * 3 \%in\% c(2,3,4)$. What do you note in the execution order of the operators?

Objects & functions

- **Objects:** identifying the values that have been *saved* in the workspace (Environment). Values are *assigned* to objects by using the <- symbol (minor and minus). To call an object, just write its name.

```
tony_32 <- 2 # assign a value to an object
tony_32 # print object
[1] 2
tony_32 <- tony_32 + 1 # update object
tony_32 # print again
[1] 3
```

- **Functions:** labels associated to sequences of commands that were programmed to return a specific output (called **value**) based on 1+ input values (called **argument**). The function name is always *followed by rounded brakes* that include the arguments. If no arguments are specified, default values are used.

```
sqrt(x = 9) # root square of the value assigned to the argument x
[1] 3
seq(from = 1, to = 5) # numerical sequence between the 'from' and 'to'
[1] 1 2 3 4 5
```

Types (classes) of objects

Logical

```
x <- TRUE  
  
x <- T # writing "TRUE" or "T" is the same  
  
class(x)  
[1] "logical"
```

Numeric

```
x <- 1.4  
  
class(x)  
[1] "numeric"
```

Integer

```
as.integer(x)  
[1] 1
```

Character

```
x <- "I like R"  
  
x # within "" the space matters!  
[1] "I like R"
```

Vector: series of values with the same class (e.g., all numeric) combined with the function `c()` (*combine*)

```
x <- c(1, 10.5, 3, 2)  
  
x + 1  
[1] 2.0 11.5 4.0 3.0  
  
sqrt(x)  
[1] 1.000000 3.240370 1.732051 1.414214  
  
y <- c("I", "like", "R")
```

Matrix: table with `nrow` * `ncol`

```
x <- matrix(1:12, nrow = 3, ncol = 4)  
rownames(x) <- y # row names  
  
x  
  
      [,1] [,2] [,3] [,4]  
I         1    4    7   10  
like      2    5    8   11  
R         3    6    9   12
```

Classes of objects: vector

An object of class *vector* is a sequence of values with the same class. It can be created with the combine function `c()` or with other functions.

```
x <- c(1, 10.5, 3, 2) # create a numeric vector
y <- 1:10 # create another numeric vector
(z <- rep(c(TRUE, FALSE), each = 2)) # create a logical vector
[1] TRUE TRUE FALSE FALSE
as.character(x) # convert the vector class from numeric to character
[1] "1" "10.5" "3" "2"
as.numeric(z) # from logical to numeric (FALSE = 0, TRUE = 1)
[1] 1 1 0 0
```

If we use a function on a vector, it applies to all the vectors' values.

```
y*2 # multiply all y values by 2
[1] 2 4 6 8 10 12 14 16 18 20
round(sqrt(y), 2) # root squared of y values, rounded at 2 digits
[1] 1.00 1.41 1.73 2.00 2.24 2.45 2.65 2.83 3.00 3.16
```

Classes of objects: vector

Some functions return a single value from a vector of values.

```
length(y) # return the number of vector elements  
[1] 10
```

For instance, those functions computing **descriptive statistics**:

```
sum(y) # sum the elements of y  
[1] 55  
max(y) # maximum value  
[1] 10  
mean(y) # mean value  
[1] 5.5  
median(y) # median value  
[1] 5.5  
var(y) # variance  
[1] 9.166667  
sd(y) # standard deviation  
[1] 3.02765
```

Classes of objects: vector

Squared brakes [] allow to **select 1+ elements** of the vector

```
tony32 <- c("one", "two", "three", "four", "five") # vector of characters
tony32[3] # select the third element
[1] "three"
tony32[3:5] # select from the third to the fifth element
[1] "three" "four" "five"
tony32[c(4, 2)] # fourth and second (not "tony32[4,2]!")
[1] "four" "two"
```

For instance, we can **select those elements meeting certain conditions** by using logical and relational operators:

```
y[y <= 3 | y > 8] # y values lower/equal to 3 or higher than 8
[1] 1 2 3 9 10
tony32[tony32 != "two"] # all values different from "two"
[1] "one" "three" "four" "five"
tony32[substr(tony32, 2, 2) == "w"] # values with the letter "w" in the 2nd position
[1] "two"
```

Classes of objects: vector

The `which()` function returns the **position** of the values meeting the condition

```
substr(tony32, 2, 2) == "w" # test equivalence for each value
[1] FALSE TRUE FALSE FALSE FALSE
which(substr(tony32, 2, 2) == "w") # position of values meeting the cond.
[1] 2
tony32[substr(tony32,2,2)== "w"] == tony32[which(substr(tony32,2,2)== "w")]
[1] TRUE
```

To **replace** 1+ values of a vector, use the symbol `<-`

```
tony32[1] <- "three"
```

To **remove** 1+ values from a vector, use the symbol `-`

```
tony32[-c(2, 4)]
[1] "three" "three" "five"
```

Classes of objects: factor

An object of class *factor* is a special type of vector used in R to work with *categorical variables* (nominal and ordinal). The possible values of a factor are called *levels*, which are ordered increasingly by default (i.e., numeric or alphabetic order).

```
as.factor(tony32) # from character to factor
[1] three two   three four   five
Levels: five four three two

# summarize the values of a vector
summary(tony32)

      Length      Class      Mode
      5 character character

# with factors, it shows the freq. by level
summary(as.factor(tony32)) # equivalent to table()

five four three   two
  1    1    2    1
```

```
(y <- rep(c(2,4,6),3)) # numeric vector
[1] 2 4 6 2 4 6 2 4 6
```

```
as.factor(y) # from numeric to factor
[1] 2 4 6 2 4 6 2 4 6
Levels: 2 4 6
```

```
summary(y) # summary of numeric vector

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
         2         2         4         4         6         6
```

```
summary(as.factor(y)) # summary of factor
2 4 6
3 3 3
```

Classes of objects: factor

Create a factor with the function `factor()`:

```
factor(x = c("C",rep("A",3),c("B","A","C")))
```

```
[1] C A A A B A C
```

```
Levels: A B C
```

```
(x <- factor(x = c("C","A","B","A"), # vector  
            levels = c("C","A","B"))) # levels
```

```
[1] C A B A
```

```
Levels: C A B
```

```
levels(x) # levels() print the levels
```

```
[1] "C" "A" "B"
```

```
factor(x, levels=c("B","A","C")) # change order
```

```
[1] C A B A
```

```
Levels: B A C
```

```
levels(x) <- c("Uno","Due","Tre") # change names
```

```
x
```

```
[1] Uno Due Tre Due
```

```
Levels: Uno Due Tre
```

Factor levels can be **ordered** by setting the argument `ordered = TRUE`, which returns an ordinal variable based on the order defined by the levels argument.

```
# unordered factor (default)
```

```
vec <- c("Maria","Mauro","Teresa","Carlo")
```

```
factor(x = vec)
```

```
[1] Maria Mauro Teresa Carlo
```

```
Levels: Carlo Maria Mauro Teresa
```

```
# ordered factor (default alphabetical order)
```

```
factor(x = vec, ordered = TRUE)
```

```
[1] Maria Mauro Teresa Carlo
```

```
Levels: Carlo < Maria < Mauro < Teresa
```

```
# ordered factor (setting a different order)
```

```
factor(x = vec, ordered = TRUE,  
      levels=c("Teresa","Mauro","Maria","Carlo"))
```

```
[1] Maria Mauro Teresa Carlo
```

```
Levels: Teresa < Mauro < Maria < Carlo
```


Classes of objects: matrix

A *matrix* is a **bidimensional structure** ($\text{nrow} * \text{ncol}$) of values with the same **class**, which can be created with the function:

```
matrix(data, nrow = , ncol = , byrow = FALSE)
```

```
(x <- matrix(1:12, nrow = 3, ncol = 4))

[,1] [,2] [,3] [,4]
[1,]  1   4   7  10
[2,]  2   5   8  11
[3,]  3   6   9  12

matrix(1:12, nrow = 3, ncol = 4, byrow = TRUE)

[,1] [,2] [,3] [,4]
[1,]  1   2   3   4
[2,]  5   6   7   8
[3,]  9  10  11  12

matrix(c("Mar", "Mau", "Ter", "Car"), nrow = 2)

[,1] [,2]
[1,] "Mar" "Ter"
[2,] "Mau" "Car"
```

To **select 1+** values from a matrix, we use again **squared brackets**, but this time we use the syntax `matrix_name[row_number, col_number]`.

```
x[1,2] # 1st row, 2nd column
[1] 4

x[2,1] # 2nd row, 1st column
[1] 2

x[1:3,2] # rows 1-3, second column
[1] 4 5 6

x[1,] # 1st row, all columns
[1] 1 4 7 10

x[,2] # 2nd column, all rows
[1] 4 5 6
```

Classes of objects: matrix

```
x # print the matrix x
```

	[,1]	[,2]	[,3]	[,4]
[1,]	1	4	7	10
[2,]	2	5	8	11
[3,]	3	6	9	12

Join 1+ matrices with `cbind()` and `rbind()`

```
cbind(x,matrix(rep(3,6),nrow=3)) # by column
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	4	7	10	3	3
[2,]	2	5	8	11	3	3
[3,]	3	6	9	12	3	3

```
rbind(x,matrix(rep(3,4),ncol=4)) # by row
```

	[,1]	[,2]	[,3]	[,4]
[1,]	1	4	7	10
[2,]	2	5	8	11
[3,]	3	6	9	12
[4,]	3	3	3	3

Setting row and column names:

```
rownames(x) <- c("a","b","c") # row names
colnames(x) <- 1:4 # column names
```

```
x
```

	1	2	3	4
a	1	4	7	10
b	2	5	8	11
c	3	6	9	12

```
c(nrow(x),ncol(x)) # No. of rows and cols = dim(x)
```

```
[1] 3 4
```

```
t(x) # transpose matrix (inverting rows and cols)
```

	a	b	c
1	1	2	3
2	4	5	6
3	7	8	9
4	10	11	12

Data structures: data.frame

A *dataframe* is a **bidimensional structure of vectors that can have different classes** (e.g., numeric, character, and factor) that can be created with the function:

```
data.frame(name_var1 = c(...), name_var2 = c(...), ...)
```

```
(x <- data.frame(Num = 1:4,  
                 Char = c("a","b","c","d"),  
                 Logi = rep(c(TRUE,FALSE),2)))
```

```
   Num Char  Logi  
1    1   a  TRUE  
2    2   b FALSE  
3    3   c  TRUE  
4    4   d FALSE  
  
str(x) # structure of the dataframe  
'data.frame':   4 obs. of  3 variables:  
 $ Num : int  1 2 3 4  
 $ Char: chr  "a" "b" "c" "d"  
 $ Logi: logi  TRUE FALSE TRUE FALSE
```

While `str(df_name)` returns the structure of a dataframe, `summary(df_name)` returns a summary for each column.

```
summary(x)
```

	Num	Char	Logi
Min.	:1.00	Length:4	Mode :logical
1st Qu.:	:1.75	Class :character	FALSE:2
Median :	:2.50	Mode :character	TRUE :2
Mean :	:2.50		
3rd Qu.:	:3.25		
Max.	:4.00		

Data structures: data.frame

Manipulating a dataframe is very similar to what we saw for matrices.

Selecting values and joining two dataframes:

```
x[2, 2:3] # 2nd row, 2nd and 3rd column
```

	Char	Logi
2	b	FALSE

```
cbind(x, data.frame(new=4:1)) # join by col
```

	Num	Char	Logi	new
1	1	a	TRUE	4
2	2	b	FALSE	3
3	3	c	TRUE	2
4	4	d	FALSE	1

```
rbind(x[1:3,], data.frame(Num=10, # join by row
                           Char="z", Logi=FALSE))
```

	Num	Char	Logi
1	1	a	TRUE
2	2	b	FALSE
3	3	c	TRUE
4	10	z	FALSE

Name and number of rows and columns:

```
rownames(x) # default = 1:nrow(x)
```

```
[1] "1" "2" "3" "4"
```

```
colnames(x)[2] # 2nd column name
```

```
[1] "Char"
```

```
nrow(x) # no. of rows
```

```
[1] 4
```

```
ncol(x) # no. of cols
```

```
[1] 3
```

Transpose of a dataframe:

```
t(x)
```

	[,1]	[,2]	[,3]	[,4]
Num	"1"	"2"	"3"	"4"
Char	"a"	"b"	"c"	"d"
Logi	"TRUE"	"FALSE"	"TRUE"	"FALSE"

Data structures: dataframe



To select a column (vector) from a dataframe, we can use the \$ symbol with the syntax `df_name$column_name`:

```
x$Char # selecting the Char column
[1] "a" "b" "c" "d"

x$Char[2] # second value from the Char column
[1] "b"

x$Char[2] == x[2,2] # equivalent commands
[1] TRUE

x$Char <- NULL # removing the Char column

x[x$Num < 3,] # selecting cases with Num < 3

  Num  Logi
1    1  TRUE
2    2 FALSE

# same result with subset(x, Num < 3)
```

Alternative way to select the columns:

```
x[, "Logi"] # write col names within ""
[1] TRUE FALSE TRUE FALSE

x[1:2, c("Num", "Logi")]

  Num  Logi
1    1  TRUE
2    2 FALSE
```

“Head” and “tail” of a dataframe:

```
head(x, n = 2) # first 2 rows

  Num  Logi
1    1  TRUE
2    2 FALSE

tail(x, 1) # last row

  Num  Logi
4    4 FALSE
```

Data structures: list

A *list* is a **collection of objects** that can have **different classes** (es. vector, matrix, and data.frame) e **different length** (contrarily to matrices and dataframes). It is the most complex and versatile structure in R, which can be created with the function

```
list(nome_oggetto1 = ..., nome_oggetto2 = ..., ...)
```

```
x <- list(Num = 1:4,
          Matr = matrix(1:12, nrow=3),
          df = x,
          lst = list(1:3,2:3))

str(x) # structure of the list

List of 4

 $ Num : int [1:4] 1 2 3 4
 $ Matr: int [1:3, 1:4] 1 2 3 4 5 6 7 8 9 10 ...
 $ df  :'data.frame':  4 obs. of  2 variables:
 ..$ Num : int [1:4] 1 2 3 4
 ..$ Logi: logi [1:4] TRUE FALSE TRUE FALSE
 $ lst :List of 2
 ..$ : int [1:3] 1 2 3
 ..$ : int [1:2] 2 3
```

To **select 1+ objects** from a list, we still use **squared brakets**

```
x[1] # single = create a sub-list
$Num
[1] 1 2 3 4

class(x[1])
[1] "list"

x[[1]] # double = extract the object
[1] 1 2 3 4

class(x[[1]])
[1] "integer"

x[[3]][2,1]
[1] 2
```

Functions & packages

Many things in R can be done by using **functions**, which always include the following elements: *function name*, *arguments* (`arg_name = arg_value` or without name, based on the default position), and returned *value*.

```
sqrt(x = c(1,2,3))  
[1] 1.000000 1.414214 1.732051  
sqrt(c(1,2,3))  
[1] 1.000000 1.414214 1.732051
```

R Help system: To know the details of any function (arguments, value, etc.), just add the `?` symbol before the function name

```
?sqrt
```

R packages: To get additional functions than those included in the base R packages, you need to install and open the related package

```
install.packages("pkg_name") # installing a package  
library(pkg_name) # open a package  
pkg_name::function_name() # using a function without loading the pkg
```

Objects, functions, and workspace

When we assign a value to an object, the latter is recorded in the **workspace**: the place that includes all objects and functions defined by the user (*Environment* tab in Rstudio).

The `ls()` function returns the name of any object and function included in the workspace, whereas the `rm()` function removes the object(s) selected:

```
x <- 1 # assign value to x
y <- 2 # assign value to y
ls() # show all objects in the workspace
[1] "x" "y"
rm(y) # remove object y
ls()
[1] "x"
```

Combining the two functions, the command `rm(list=ls())` **empties the workspace**, removing all objects and functions (very **useful at the beginning of any script!**)

```
rm(list = ls()) # remove all objects & functions
```


Objects & functions from inside R:

Default packages

Some packages are **installed in R by default**, they do not require to be opened, and they are not shown in the workspace.

```
rm(list=ls()) # remove all objects & functions
head(sleep,4) # 'sleep' dataset from the 'datasets' pkg
  extra group ID
1   0.7      1  1
2  -1.6      1  2
3  -0.2      1  3
4  -1.2      1  4
mean(sleep$extra) # 'mean()' function from the 'base' pkg
[1] 1.54
letters[2] # constant values from the 'base' pkg
[1] "b"
ls() # yet the workspace looks empty!
character(0)
```

Objects & functions from outside R:

The working directory

To read a file from a specific folder, you should first get (and possibly set) the **working directory**, that is the folder where input files are searched and where output files are saved.

```
getwd() # returns the current working directory
[1] "C:/Users/mengh/OneDrive/Desktop/PHD/Didattica/advancedDataAnalysisPsy"
dir()[1:3] # names of the first 3 files in the WD
[1] "~$img.pptx"
[2] "~$llabus corso ADVANCED DATA ANALYSIS FOR PSYCHOLOGICAL SCIENCE.docx"
[3] "1-course-intro.Rmd"
```

```
setwd("data") # moving WD to the 'data' subfolder
setwd("C:/Users/mengh/OneDrive/Desktop") # set new WD
```

Trick with RStudio: when you start a new project (e.g., data analysis for the thesis, report), create a new **R project** (.Rproj) from the menu **File > New R Project** by selecting an existing folder or by creating a new one, which will be the default WD for any file related to that project.

Reading & saving a dataset (1/2)

The first step in any data analysis with R is to read a dataset and save it in the workspace. To do so, we need to use a specific function based on the data format (e.g., CSV, xlsx, txt, svy). The default R data format is RData.

```
# file .RData  
load(file = "data/studqs.RData") # import  
save(qs, file = "data/studqs.RData") # export
```

R can read and export further data formats, some of which require to install additional packages.

```
# file .CSV (comma separated values)  
qs <- read.csv(file = "data/studqs.csv") # import  
write.csv(x = qs, "data/studqs.csv", row.names = FALSE) # export  
  
# file .SAV (da SPSS)  
library(foreign)  
qs <- read.spss("data/studqs.sav", to.data.frame=TRUE) # import
```

Reading & saving a dataset (2/2)

A fast (but poorly reproducible!) way to read a file in R is by using the function `file.choose()` rather than the name of the file. This allows reading a file without setting the WD.

```
qs <- read.csv(file = file.choose())
```

If you have doubts on the function to be used for reading a file, you can use the Rstudio menu **File > Import Dataset**

Hands on :

- Open or create a file `.xlsx` on your PC, save it in a `.csv` (comma separated values) format and read it in R
- Now try to directly read the `.xlsx` file (if you have doubts, just Google “*how to read xlsx file with R*” 😊)
- In both cases, note the class and the structure of the imported object

Hands on : Student questionnaire

1. Download the files `questionarioStudenti.RData` and `questionarioStudenti.csv` from [Github](#) (data folder; select **File** > **Raw** > **Download** or right click > **Save as**) or [Moodle](#) (data folder), save the file in a folder and set that folder as the working directory.
2. Read both files with R: what are their class? What is the class of the included variables?
3. Use the `describe()` function from the `psych` package to compute the descriptive statistics of the variable `numVar` and use the function `hist()` to visualize its histogram plot (try changing the argument `breaks`).
4. Use the function `table()` to create a frequency table of the variable `Q02`
5. Repeat point #5 but only consider the students that replied “Si” (yes) to the item `Q01`
6. Cross the frequency of the variables `Q02` and `Q03` using `table()`

Descriptive statistics (univariate)

```
x <- c(1,1,1,2,8,9) # create numeric vector
```

```
c(mean(x),median(x)) # mean & median
```

```
[1] 3.666667 1.500000
```

```
as.numeric(which.max(table(x))) # mode
```

```
[1] 1
```

```
c(var(x),sd(x)) # variance & standard dev.
```

```
[1] 14.266667 3.777124
```

```
quantile(x,probs=0.90) # 90o percentile
```

```
90%
```

```
8.5
```

```
quantile(x,probs=c(0.25,0.50,0.75,1)) # quartiles
```

```
25% 50% 75% 100%
```

```
1.0 1.5 6.5 9.0
```

```
round(rank(x)/length(x),2) # sample rank
```

```
[1] 0.33 0.33 0.33 0.67 0.83 1.00
```

```
table(x) # absolute frequencies
```

```
x
```

```
1 2 8 9
```

```
3 1 1 1
```

```
round(table(x)/length(x),2) # relative freq
```

```
x
```

```
1 2 8 9
```

```
0.50 0.17 0.17 0.17
```

```
cumsum(x) # cumulative sum
```

```
[1] 1 2 3 5 13 22
```

```
cumsum(table(x)) # absolute cumulative sums
```

```
1 2 8 9
```

```
3 4 5 6
```

```
round(cumsum(table(x)/length(x),2)
```

```
1 2 8 9
```

```
0.50 0.67 0.83 1.00
```

Descriptive statistics (bivariate)

```
y <- -x - 1 # values inversely proport. to x
z <- round(rnorm(n=length(x)),1) # random values
(df <- data.frame(x,y,z)) # new data frame
```

	x	y	z
1	1	-2	-0.5
2	1	-2	-0.4
3	1	-2	0.4
4	2	-3	-0.6
5	8	-9	-0.7
6	9	-10	0.8

Correlation & covariance

```
cov(x,y) # covariance btw x & y
[1] -14.26667

cor(x,y) # correlation
[1] -1
```

,

With more than 2 variables

```
cov(df) # covariance matrix
```

	x	y	z
x	14.2666667	-14.2666667	0.7133333
y	-14.2666667	14.2666667	-0.7133333
z	0.7133333	-0.7133333	0.3786667

```
cor(df) # correlation matrix
```

	x	y	z
x	1.0000000	-1.0000000	0.3069041
y	-1.0000000	1.0000000	-0.3069041
z	0.3069041	-0.3069041	1.0000000

Graphics in R: Main functions

High-level functions (stand alone)

basic function (plot depends on object class)

`plot()` *# scatter plot*

distributions

`boxplot()` *# boxplot*

`qqnorm()` *# quantile-quantile plot*

frequencies

`barplot()` *# barplot (categorical variables)*

`hist()` *# histogram (continuous variables)*

`pie()` *# pie chart*

interactions

`interaction.plot()`

Lower-level functions

(add elements to the former)

`points()` *# add dots*

`lines()` *# add lines*

`text()` *# add text*

linear regression line

`abline(a = ..., b = ...)`

more complex elements

`rect()` *# add rectangles*

`polygon()` *# add polygons*

other graphical features

`axis()` *# change the plot axes*

`legend()` *# add a legend*

Graphical parameters

Run the command `?par` to see all graphical parameters that can be changed.

Some of them can be set by using the arguments of graphical functions:

```
# dimension

cex = 2 # text & simbol dimension (multiplier)

lwd = 0.5 # line width


# color & shape

col = "red" # (see ?colors)

lty = 2 # type of line (1=solid, 2=dashed, ...)

pch = 19 # shape of dots (see ?points)


# titles

main = "Plot title"

xlab = "x-axis title"
```

Other parameters should be set within the `par()` function, which should be executed before generating the plot:

```
# margins (bottom,left,top,right)

mai # in inches

mar # in text lines


# multiple panels

mfrow=c(nrow, ncol)

# e.g. two plots next to each other

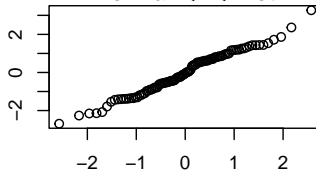
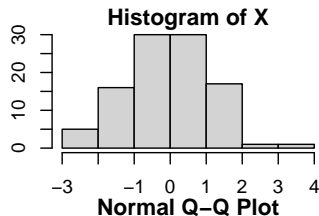
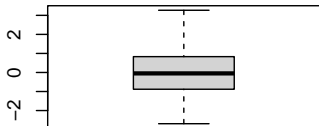
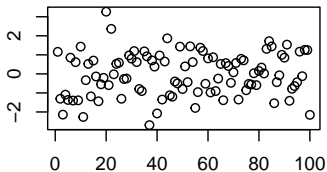
mfrow=c(2,1)

# es. one above and one below

mfrow=c(1,2)
```

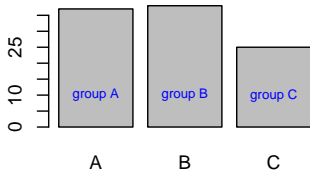
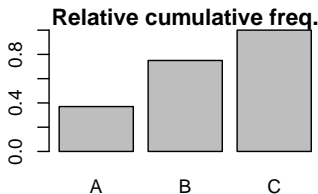
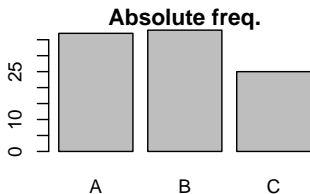
Univariate distributions: Continuous variables

```
X <- rnorm(n=100, mean=0, sd=1) # generating random values from the stand. norm. distr.  
par(mfrow = c(2,2), mai = c(0.3,0.5,0.2,0.5)) # plot 4 graphs in a single window  
plot(X) # scatter plot (variability)  
hist(X) # histogram: frequency distr. by class  
boxplot(X) # box plot: (1st & 3rd quartile +/- 1.5 IQR)  
qqnorm(X) # Q-Q plot: cumulative distribution of X vs. normal cumulative distr.
```



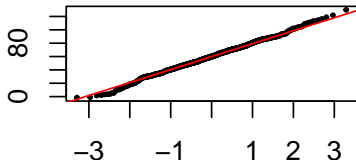
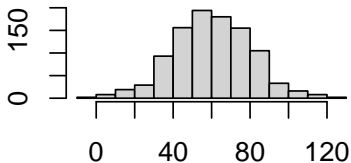
Univariate distributions: Categorical variables

```
Y <- factor(sample(x=c("A","B","C"),size=100,replace=TRUE)) # generating random factor
par(mfrow = c(2,2), mai = c(0.3,0.5,0.2,0.5)) # plot 4 graphs in a single window
barplot(table(Y), main="Absolute freq.") # bar plot: absolute frequencies
barplot(round(cumsum(table(Y)/length(Y)),2), main="Relative cumulative freq.")
barplot(table(Y)) # same graph but adding blue text
text(x=c(0.7,1.9,3.1),y=10,labels=c("group A","group B","group C"),col="blue",cex=0.7)
```

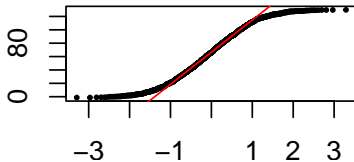
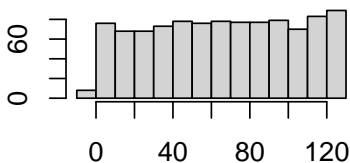


Assessing the normality of a distribution (1/3)

```
X <- rnorm(n = 1000, mean = 60, sd = 20) # norm
par(mfrow=c(2,1),mai=rep(0.3,4),mar=rep(1.8,4))
hist(X,main="",xlab="",ylab="") # histogram
qqnorm(X,cex=0.5,main="",pch=20) # Q-Q plot
qqline(X,col="red") # adding normal line
```



```
Y <- runif(n=1000,min=min(X),max=max(X)) # unif
par(mfrow=c(2,1),mai=rep(0.3,4),mar=rep(1.8,4))
hist(Y,main="",xlab="",ylab="") # histogram
qqnorm(Y,cex=0.5,main="",pch=20) # Q-Q plot
qqline(Y,col="red") # adding normal line
```



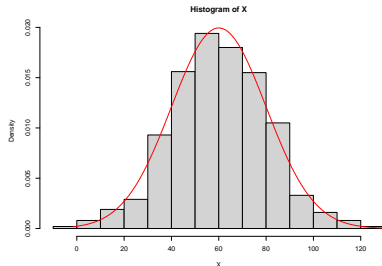
Assessing the normality of a distribution (2/3)

```
# Empirical distribution
```

```
hist(X, freq = FALSE)
```

```
# Theoretical (expected) distribution
```

```
curve(dnorm(x, mean = 60, sd = 20),  
      from = min(X), to = max(X),  
      col = "red", lwd = 2,  
      add = TRUE)
```

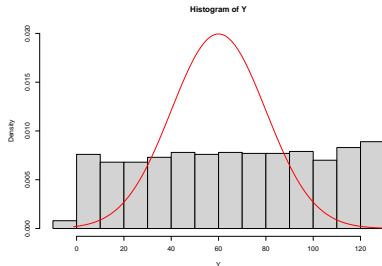


```
# Empirical distribution
```

```
hist(Y, freq = FALSE, ylim = c(0,0.02))
```

```
# Theoretical (expected) distribution
```

```
curve(dnorm(x, mean = 60, sd = 20),  
      from = min(Y), to = max(Y),  
      col = "red", lwd = 2,  
      add = TRUE)
```



Assessing the normality of a distribution (3/3)

```
library(moments) # skewness & kurtosis
c(skewness(X), kurtosis(X)) # (ok if ~ 0)
[1] -0.02306247  3.19064442
# Kolmogorov-Smirnov test (H0: X is normal)
ks.test(X, y="pnorm", mean=mean(X), sd=sd(X))
```

Asymptotic one-sample Kolmogorov-Smirnov test

```
data: X
D = 0.014883, p-value = 0.9797
alternative hypothesis: two-sided
# Shapiro-Wilk test (H0: X is normal)
shapiro.test(X)
```

Shapiro-Wilk normality test

```
data: X
W = 0.99794, p-value = 0.2599
```

```
# skewness & kurtosis
c(skewness(Y), kurtosis(Y)) # (ok if ~ 0)
[1] -0.03894971  1.82192786
# Kolmogorov-Smirnov test (H0: X is normal)
ks.test(Y, y="pnorm", mean=mean(Y), sd=sd(Y))
```

Asymptotic one-sample Kolmogorov-Smirnov test

```
data: Y
D = 0.058807, p-value = 0.001983
alternative hypothesis: two-sided
# Shapiro-Wilk test (H0: X is normal)
shapiro.test(Y)
```

Shapiro-Wilk normality test

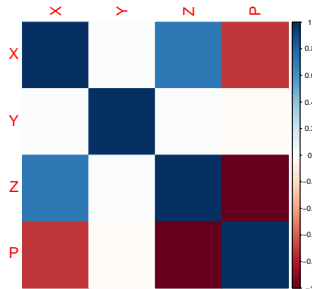
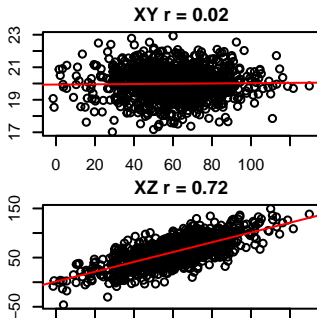
```
data: Y
W = 0.95518, p-value < 2.2e-16
```

Bivariate distributions: Continuous variables

```
Y <- rnorm(1000, mean=20, sd=1) # Y (independent from X)
Z <- X + rnorm(1000, sd = 20) # Z (proportional to X)
# Scatter plot
par(mfrow=c(2,1),mai=rep(0.3,4),mar=rep(1.8,4),cex=.5)
plot(X, Y, main = paste("XY r =",round(cor(X, Y),2)))
abline(lm(Y~X),col="red") # linear regression line
plot(X, Z, main = paste("XZ r =",round(cor(X, Z),2)))
abline(lm(Z~X),col="red") # tilde (~) with Alt + 1-2-6
```

```
df <- data.frame(X, Y, Z, P = -Z)

# Correlation plot (or heat plot)
library(corrplot)
corrplot(cor(df),method="color",
         tl.cex=2,cl.cex=1,
         mar = c(0,0,0,15))
```



Bivariate distributions: Categorical variables

```
# create data.frame with sex (M/F) & education level (primary, middle, highsch, univ.)  
df <- data.frame(sex = sample(x = rep(c("F","M"),50), size = 100),  
                 edu = sample(x = rep(c("prim","mid","high","uni"),25), size = 100))
```

```
table(df) # crossed frequencies
```

```
edu
```

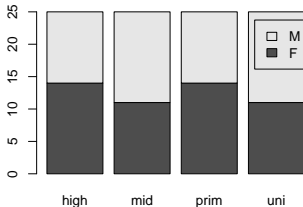
```
sex high mid prim uni
```

```
F 14 11 14 11
```

```
M 11 14 11 14
```

```
par(mai=rep(0.4,4))
```

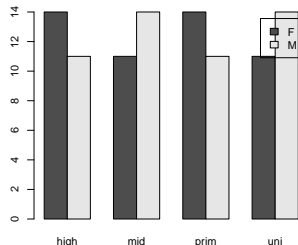
```
barplot(table(df),legend.text = TRUE) # bar plot
```



```
# version with 'sided groups'
```

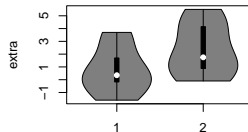
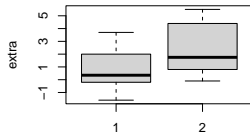
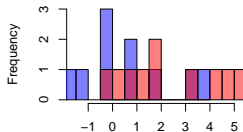
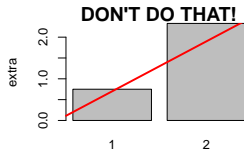
```
par(mai=rep(0.4,4))
```

```
barplot(table(df),legend.text = TRUE,  
        beside = TRUE)
```



Bivariate distributions: Continuous & categorical

```
par(mfrow = c(2,2), mai = c(0.5,1,0.2,0.5)) # using the sleep dataset (included in R base)
barplot(tapply(sleep$extra, sleep$group, FUN=mean), # barplot for means (DON'T DO THAT!)
        main="DON'T DO THAT!", cex.main=1.5, ylab = "extra"); abline(a=0,b=1,col="red",lwd=2)
hist(sleep[sleep$group==1,"extra"], breaks=10, col = rgb(0,0,1,alpha=0.5),
     xlim = c(min(sleep$extra), max(sleep$extra)), xlab="", main="") # hist by group
hist(sleep[sleep$group==2,"extra"], breaks=10, col = rgb(1,0,0,alpha=0.5), add = TRUE)
boxplot(extra ~ group, data = sleep) # box plot by group
library(vioplot); violoplot(extra ~ group, data = sleep) # violin plot: density by group
```



Saving a plot

```
# 1. open new file with PNG format

png(filename = "grafico.png", # file name
      width = 350, height = 350) # width & height

# 2. create the plot
plot(sleep$group, sleep$extra,
      xlab = "group", ylab = "extra", col = "#2E9FDF")

# 3. close the file
dev.off() # this will save the plot in your WD

pdf

2
```

Alternative formats (see ?png)

png() *# Portable Network Graphics*

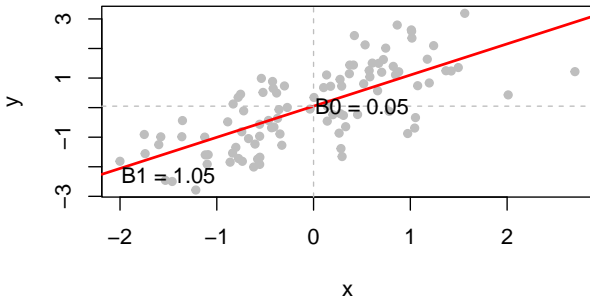
bpm() *# bitmap*

jpeg() *# Joint Photographic Experts Group*

tiff() *# Tagged Image File Format*

Linear regression

- **Regression** allows to determinate the asymmetrical and functional relation between two variables, to quantify whether one variable (**predictor**) predicts another variable (**dependent** or **response**)
- **Linear regression** determines the link between two variables through a linear function: $y_i = \beta_0 + \beta_1 x_i + \epsilon_i$
- Such a function can be graphically represented as a **straight line**, where β_0 is the **intercept** (value assumed by y when x = 0) and β_1 is the **slope** (predicted change in y when x increases by 1 unit)



Fitting a linear regression model (1/3)

Using the dataset **gambling** from the **ADati** pkg. We want to fit models that predict **frequency** of gamig behavior based on **gender** and **risk** (risk perception).

```
data(gambling, package = "ADati")
```

R uses the `lm()` function with the arguments **formula** (syntax: **dependent** ~ **pred1** + **pred2** + ..., where the symbol “~” is the **tilde**, on Windows: **Alt + 126**) and **data** (dataframe that includes the variables in the formula)

- **Null model:** **frequency** scores are only predicted by the **intercept** β_0 , that is the mean value of **frequency** in the sample.

```
m0 <- lm(formula = frequency ~ 1, data = gambling)
```

```
coefficients(m0) # coefficient estimates (intercept)
```

```
(Intercept)
```

```
10.55997
```

```
mean(gambling$frequency) # mean of the frequency variable
```

```
[1] 10.55997
```

Fitting a linear regression model (2/3)

- **Simple linear regression model:** frequency is predicted by the intercept β_0 (expected value of frequency when gender is "f") and the slope coefficient β_1 , expressing the 'mean' difference $m - f$.

```
m1 <- lm(formula = frequency ~ gender, data = gambling)
```

```
coefficients(m1)
```

```
(Intercept)      genderm  
  9.336685      3.042012
```

```
tapply(gambling$frequency, gambling$gender, mean) # mean of frequency for f & m  
      f      m  
9.336685 12.378697
```

Fitting a linear regression model (3/3)

- **Multiple linear regression model:** frequency is predicted by the intercept β_0 (expected frequency when `gender = f` and `risk = 0`) and the coefficients β_1 ('mean' gender difference when `risk = 0`) and β_2 (effect of `risk` net to gender differences)

```
m2 <- lm(formula = frequency ~ risk + gender, data = gambling)
```

```
coefficients(m2)
```

(Intercept)	risk	genderm
18.951650	-1.393899	2.008555

- **Interactive model:** frequency is predicted by the intercept β_0 and the coefficients β_1 ('media' gender difference when `risk = 0`) and β_2 (effect of `risk` when `gender = "f"`) and β_3 (effect of `risk` when `gender = "m"`)

```
m3 <- lm(formula = frequency ~ risk * gender, data = gambling)
```

```
coefficients(m3)
```

(Intercept)	risk	genderm	risk:genderm
16.9206907	-1.0994674	5.9744077	-0.6087177

Assessing linear model assumptions (1/3)

1. Independence of observations: all pairs of errors ϵ_i and ϵ_j are independent for any $i \neq j$. This assumption does not require particular procedures, just think: are my observations independent?

2. Independence of predictors and errors: residuals (estimates of the error component ϵ_i) should not be associated with the predictors **gender** and **risk**.

But what are residuals?

```
gambling$RESIDUALS <- residuals(m3) # model residuals with residuals()
gambling$PREDICTED <- fitted(m3) # predicted values with fitted()
head(gambling[,c("frequency", "PREDICTED", "RESIDUALS")], 5) # RESIDUALS = OBSERVED - PREDICTED
```

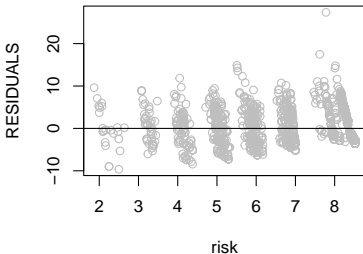
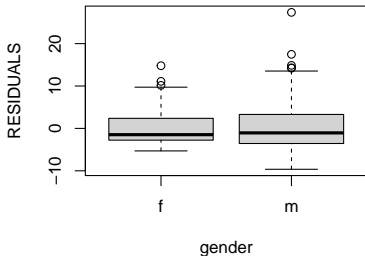
	frequency	PREDICTED	RESIDUALS
6	4.78	7.553228	-2.7732283
8	36.98	9.605418	27.3745820
10	4.78	7.553228	-2.7732283
18	11.16	11.335396	-0.1753962
19	5.62	7.828095	-2.2080952

Assessing linear model assumptions (2/3)

2. Independence of predictors and errors: residuals (estimates of the error component ϵ_i) should not be associated with the predictors **gender** and **risk**. Here, we do not see substantial gender differences or a substantial relationship with **risk**.

3. Constant variance (Homogeneity of variances): residual variance is constant for any value of X . Here, we see that residual variance is slightly (but not substantially) higher in males than in females.

```
par(mfrow=c(1,2),mar=c(5,4,0,2))+0.1)
plot(RESIDUALS ~ gender, data=gambling)
plot(RESIDUALS ~ risk, data=gambling, col="gray")
abline(lm(RESIDUALS ~ risk, data=gambling))
```



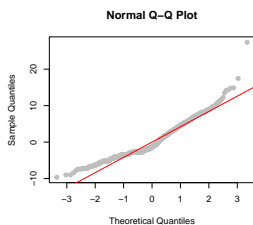
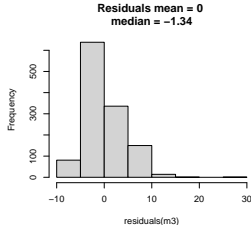
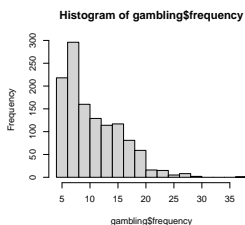
Assessing linear model assumptions (3/3)

4. Linearity: the expected value of residuals for a given X value is zero. Practically, residuals should have mean zero (as in this case).

5. Normality: residuals should be normally distributed.

Here, the fit is not very nice but we can accept it...

```
par(mfrow=c(1,3))  
hist(gambling$frequency) # if residuals are normal, Y is normal as well  
hist(residuals(m3),main=paste("Residuals mean =",round(mean(residuals(m3)),2),  
                             "\nmedian =",round(median(residuals(m3)),2)))  
qqnorm(residuals(m3), col="gray", pch=19)  
qqline(residuals(m3), col="red")
```



Comparing multiple models

Likelihood ratio test: testing the ratio of the loglikelihood (probability of data functional to the parameter estimates)

```
library(lmtest)
lrtest(m0,m1,m2,m3)

Likelihood ratio test

Model 1: frequency ~ 1
Model 2: frequency ~ gender
Model 3: frequency ~ risk + gender
Model 4: frequency ~ risk * gender

#Df  LogLik Df  Chisq Pr(>Chisq)
1    2 -3687.3
2    3 -3629.4  1 115.843 < 2.2e-16 ***
3    4 -3480.6  1 297.484 < 2.2e-16 ***
4    5 -3472.6  1  16.149 5.853e-05 ***

---

Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Information criteria: penalizing models based on specific criteria (e.g., low likelihood and parsimony)

```
# Akaike information criterion (AIC)
AIC(m0,m1,m2,m3)$AIC # the lower the better
[1] 7378.599 7264.756 6969.272 6955.122

# Akaike weights: from 0 (-) to 1 (+)
library(MuMIn)
Weights(AIC(m0,m1,m2,m3)) # Aw
model weights
[1] 0.000 0.000 0.001 0.999
```

Inference on regression coefficients

Based on the NHST approach, we can test the statistical significance of the regression coefficients. In R, such tests are conducted by default when we fit the model.

```
summary(m3)$coefficients # command summary() shows the model summary
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	16.9206907	0.7416563	22.814733	3.337032e-96
risk	-1.0994674	0.1051711	-10.454087	1.484438e-24
genderm	5.9744077	1.0163116	5.878519	5.339754e-09
risk:genderm	-0.6087177	0.1512209	-4.025355	6.040343e-05

Determination coefficient R^2 :

```
# the model explains the 29% of the variance in frequency
```

```
summary(m3)$r.squared
```

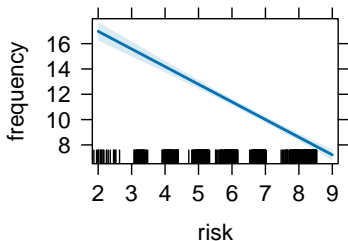
```
[1] 0.2965382
```

Visualizing the coefficient estimates

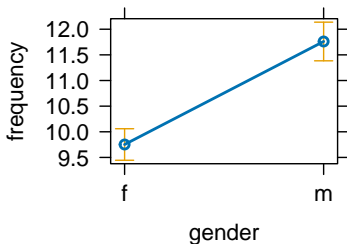
To visualize the estimated “effects” (regression line + 95% confidence intervals), we can use the `effects` package.

```
library(effects)
plot(allEffects(m2))
```

risk effect plot



gender effect plot



Other resources: First steps with R

- R Core Team. *The R Manuals*. Manuals in pdf format:
<https://cran.r-project.org/manuals.html> (particularly: “*An Introduction to R (with many examples)*”; “*R Data Import/Export*”)
- UniGlasgow - PsyTheacR. Interactive and free course on data analysis with R:
<https://psyteachr.github.io/>
- Dalgaard, P. (2008). *Introductory statistics with R*. New York: Springer.
- Murrell (2011) *R Graphics* (II edition). New York: Chapman and Hall/CRC. Examples with code from the website:
<https://www.stat.auckland.ac.nz/~paul/RGraphics/rgraphics.html>